

Programming Language—Common Lisp

18. Hash Tables

Version 15.17R, X3J13/94-101R.
Fri 12-Aug-1994 6:35pm EDT

18.1 Hash Table Concepts

18.1.1 Hash-Table Operations

Figure 18–1 lists some *defined names* that are applicable to *hash tables*. The following rules apply to *hash tables*.

- A *hash table* can only associate one value with a given key. If an attempt is made to add a second value for a given key, the second value will replace the first. Thus, adding a value to a *hash table* is a destructive operation; the *hash table* is modified.
- There are four kinds of *hash tables*: those whose keys are compared with **eq**, those whose keys are compared with **eql**, those whose keys are compared with **equal**, and those whose keys are compared with **equalp**.
- *Hash tables* are created by **make-hash-table**. **gethash** is used to look up a key and find the associated value. New entries are added to *hash tables* using **setf** with **gethash**. **remhash** is used to remove an entry. For example:

```
(setq a (make-hash-table)) → #<HASH-TABLE EQL 0/120 32536573>
(setf (gethash 'color a) 'brown) → BROWN
(setf (gethash 'name a) 'fred) → FRED
(gethash 'color a) → BROWN, true
(gethash 'name a) → FRED, true
(gethash 'pointy a) → NIL, false
```

In this example, the symbols **color** and **name** are being used as keys, and the symbols **brown** and **fred** are being used as the associated values. The *hash table* has two items in it, one of which associates from **color** to **brown**, and the other of which associates from **name** to **fred**.

- A key or a value may be any *object*.
- The existence of an entry in the *hash table* can be determined from the *secondary value* returned by **gethash**.

clrhash	hash-table-p	remhash
gethash	make-hash-table	sxhash
hash-table-count	maphash	

Figure 18–1. Hash-table defined names

18.1.2 Modifying Hash Table Keys

The function supplied as the `:test` argument to **make-hash-table** specifies the ‘equivalence test’ for the *hash table* it creates.

An *object* is ‘visibly modified’ with regard to an equivalence test if there exists some set of *objects* (or potential *objects*) which are equivalent to the *object* before the modification but are no longer equivalent afterwards.

If an *object* O_1 is used as a key in a *hash table* H and is then visibly modified with regard to the equivalence test of H , then the consequences are unspecified if O_1 , or any *object* O_2 equivalent to O_1 under the equivalence test (either before or after the modification), is used as a key in further operations on H . The consequences of using O_1 as a key are unspecified even if O_1 is visibly modified and then later modified again in such a way as to undo the visible modification.

Following are specifications of the modifications which are visible to the equivalence tests which must be supported by *hash tables*. The modifications are described in terms of modification of components, and are defined recursively. Visible modifications of components of the *object* are visible modifications of the *object*.

18.1.2.1 Visible Modification of Objects with respect to EQ and EQL

No *standardized function* is provided that is capable of visibly modifying an *object* with regard to `eq` or `eql`.

18.1.2.2 Visible Modification of Objects with respect to EQUAL

As a consequence of the behavior for **equal**, the rules for visible modification of *objects* not explicitly mentioned in this section are inherited from those in Section 18.1.2.1 (Visible Modification of Objects with respect to EQ and EQL).

18.1.2.2.1 Visible Modification of Conses with respect to EQUAL

Any visible change to the *car* or the *cdr* of a *cons* is considered a visible modification with regard to **equal**.

18.1.2.2.2 Visible Modification of Bit Vectors and Strings with respect to EQUAL

For a *vector* of *type* **bit-vector** or of *type* **string**, any visible change to an *active element* of the *vector*, or to the *length* of the *vector* (if it is *actually adjustable* or has a *fill pointer*) is considered a visible modification with regard to **equal**.

18.1.2.3 Visible Modification of Objects with respect to EQUALP

As a consequence of the behavior for **equalp**, the rules for visible modification of *objects* not explicitly mentioned in this section are inherited from those in Section 18.1.2.2 (Visible Modification of Objects with respect to EQUAL).

18.1.2.3.1 Visible Modification of Structures with respect to EQUALP

Any visible change to a *slot* of a *structure* is considered a visible modification with regard to **equalp**.

18.1.2.3.2 Visible Modification of Arrays with respect to EQUALP

In an *array*, any visible change to an *active element*, to the *fill pointer* (if the *array* can and does have one), or to the *dimensions* (if the *array* is *actually adjustable*) is considered a visible modification with regard to **equalp**.

18.1.2.3.3 Visible Modification of Hash Tables with respect to EQUALP

In a *hash table*, any visible change to the count of entries in the *hash table*, to the keys, or to the values associated with the keys is considered a visible modification with regard to **equalp**.

Note that the visibility of modifications to the keys depends on the equivalence test of the *hash table*, not on the specification of **equalp**.

18.1.2.4 Visible Modifications by Language Extensions

Implementations that extend the language by providing additional mutator functions (or additional behavior for existing mutator functions) must document how the use of these extensions interacts with equivalence tests and *hash table* searches.

Implementations that extend the language by defining additional acceptable equivalence tests for *hash tables* (allowing additional values for the **:test** argument to **make-hash-table**) must document the visible components of these tests.

hash-table

System Class

Class Precedence List:

hash-table, t

Description:

Hash tables provide a way of mapping any *object* (a *key*) to an associated *object* (a *value*).

See Also:

Section 18.1 (Hash Table Concepts), Section 22.1.3.13 (Printing Other Objects)

Notes:

The intent is that this mapping be implemented by a hashing mechanism, such as that described in Section 6.4 “Hashing” of *The Art of Computer Programming, Volume 3* (pp506-549). In spite of this intent, no *conforming implementation* is required to use any particular technique to implement the mapping.

make-hash-table

Function

Syntax:

make-hash-table &key *test size rehash-size rehash-threshold* → *hash-table*

Arguments and Values:

test—a *designator* for one of the *functions* **eq**, **eq1**, **equal**, or **equalp**. The default is **eq1**.

size—a non-negative *integer*. The default is *implementation-dependent*.

rehash-size—a *real* of *type* (or (integer 1 *) (float (1.0) *)). The default is *implementation-dependent*.

rehash-threshold—a *real* of *type* (real 0 1). The default is *implementation-dependent*.

hash-table—a *hash table*.

Description:

Creates and returns a new *hash table*.

test determines how *keys* are compared. An *object* is said to be present in the *hash-table* if that *object* is the *same* under the *test* as the *key* for some entry in the *hash-table*.

size is a hint to the *implementation* about how much initial space to allocate in the *hash-table*. This information, taken together with the *rehash-threshold*, controls the approximate number of entries which it should be possible to insert before the table has to grow. The actual size might be

rounded up from *size* to the next ‘good’ size; for example, some *implementations* might round to the next prime number.

rehash-size specifies a minimum amount to increase the size of the *hash-table* when it becomes full enough to require rehashing; see *rehash-threshold* below. If *rehash-size* is an *integer*, the expected growth rate for the table is additive and the *integer* is the number of entries to add; if it is a *float*, the expected growth rate for the table is multiplicative and the *float* is the ratio of the new size to the old size. As with *size*, the actual size of the increase might be rounded up.

rehash-threshold specifies how full the *hash-table* can get before it must grow. It specifies the maximum desired hash-table occupancy level.

The *values* of *rehash-size* and *rehash-threshold* do not constrain the *implementation* to use any particular method for computing when and by how much the size of *hash-table* should be enlarged. Such decisions are *implementation-dependent*, and these *values* only hints from the *programmer* to the *implementation*, and the *implementation* is permitted to ignore them.

Examples:

```
(setq table (make-hash-table)) → #<HASH-TABLE EQL 0/120 46142754>
(setf (gethash "one" table) 1) → 1
(gethash "one" table) → NIL, false
(setq table (make-hash-table :test 'equal)) → #<HASH-TABLE EQUAL 0/139 46145547>
(setf (gethash "one" table) 1) → 1
(gethash "one" table) → 1, T
(make-hash-table :rehash-size 1.5 :rehash-threshold 0.7)
→ #<HASH-TABLE EQL 0/120 46156620>
```

See Also:

gethash, hash-table

hash-table-p

Function

Syntax:

hash-table-p *object* → *generalized-boolean*

Arguments and Values:

object—an *object*.

generalized-boolean—a *generalized boolean*.

Description:

Returns *true* if *object* is of type **hash-table**; otherwise, returns *false*.

Examples:

```
(setq table (make-hash-table)) → #<HASH-TABLE EQL 0/120 32511220>
(hash-table-p table) → true
(hash-table-p 37) → false
(hash-table-p '((a . 1) (b . 2))) → false
```

Notes:

```
(hash-table-p object) ≡ (typep object 'hash-table)
```

hash-table-count

Function

Syntax:

```
hash-table-count hash-table → count
```

Arguments and Values:

hash-table—a *hash table*.

count—a non-negative *integer*.

Description:

Returns the number of entries in the *hash-table*. If *hash-table* has just been created or newly cleared (see `clrhash`) the entry count is 0.

Examples:

```
(setq table (make-hash-table)) → #<HASH-TABLE EQL 0/120 32115135>
(hash-table-count table) → 0
(setf (gethash 57 table) "fifty-seven") → "fifty-seven"
(hash-table-count table) → 1
(dotimes (i 100) (setf (gethash i table) i)) → NIL
(hash-table-count table) → 100
```

Affected By:

`clrhash`, `remhash`, `setf` of `gethash`

See Also:

`hash-table-size`

Notes:

The following relationships are functionally correct, although in practice using **hash-table-count** is probably much faster:

```
(hash-table-count table) ≡  
(loop for value being the hash-values of table count t) ≡  
(let ((total 0))  
  (maphash #'(lambda (key value)  
                (declare (ignore key value))  
                (incf total))  
    table)  
  total)
```

hash-table-rehash-size

Function

Syntax:

hash-table-rehash-size *hash-table* → *rehash-size*

Arguments and Values:

hash-table—a *hash table*.

rehash-size—a *real* of *type* (or (integer 1 *) (float (1.0) *)).

Description:

Returns the current rehash size of *hash-table*, suitable for use in a call to **make-hash-table** in order to produce a *hash table* with state corresponding to the current state of the *hash-table*.

Examples:

```
(setq table (make-hash-table :size 100 :rehash-size 1.4))  
→ #<HASH-TABLE EQL 0/100 2556371>  
(hash-table-rehash-size table) → 1.4
```

Exceptional Situations:

Should signal an error of *type* **type-error** if *hash-table* is not a *hash table*.

See Also:

make-hash-table, **hash-table-rehash-threshold**

Notes:

If the hash table was created with an *integer* rehash size, the result is an *integer*, indicating that the rate of growth of the *hash-table* when rehashed is intended to be additive; otherwise, the result

is a *float*, indicating that the rate of growth of the *hash-table* when rehashed is intended to be multiplicative. However, this value is only advice to the *implementation*; the actual amount by which the *hash-table* will grow upon rehash is *implementation-dependent*.

hash-table-rehash-threshold

Function

Syntax:

`hash-table-rehash-threshold hash-table → rehash-threshold`

Arguments and Values:

hash-table—a *hash table*.

rehash-threshold—a *real* of *type* `(real 0 1)`.

Description:

Returns the current rehash threshold of *hash-table*, which is suitable for use in a call to **make-hash-table** in order to produce a *hash table* with state corresponding to the current state of the *hash-table*.

Examples:

```
(setq table (make-hash-table :size 100 :rehash-threshold 0.5))  
→ #<HASH-TABLE EQL 0/100 2562446>  
(hash-table-rehash-threshold table) → 0.5
```

Exceptional Situations:

Should signal an error of *type* **type-error** if *hash-table* is not a *hash table*.

See Also:

`make-hash-table`, `hash-table-rehash-size`

hash-table-size

Function

Syntax:

`hash-table-size hash-table` \rightarrow *size*

Arguments and Values:

hash-table—a *hash table*.

size—a non-negative *integer*.

Description:

Returns the current size of *hash-table*, which is suitable for use in a call to **make-hash-table** in order to produce a *hash table* with state corresponding to the current state of the *hash-table*.

Exceptional Situations:

Should signal an error of *type* **type-error** if *hash-table* is not a *hash table*.

See Also:

`hash-table-count`, `make-hash-table`

hash-table-test

Function

Syntax:

`hash-table-test hash-table` \rightarrow *test*

Arguments and Values:

hash-table—a *hash table*.

test—a *function designator*. For the four *standardized hash table test functions* (see **make-hash-table**), the *test* value returned is always a *symbol*. If an *implementation* permits additional tests, it is *implementation-dependent* whether such tests are returned as *function objects* or *function names*.

Description:

Returns the test used for comparing *keys* in *hash-table*.

Exceptional Situations:

Should signal an error of *type* **type-error** if *hash-table* is not a *hash table*.

See Also:

`make-hash-table`

gethash

Accessor

Syntax:

`gethash key hash-table &optional default` → *value*, *present-p*
(`setf (gethash key hash-table &optional default) new-value`)

Arguments and Values:

key—an *object*.

hash-table—a *hash table*.

default—an *object*. The default is `nil`.

value—an *object*.

present-p—a *generalized boolean*.

Description:

Value is the *object* in *hash-table* whose *key* is the *same* as *key* under the *hash-table*'s equivalence test. If there is no such entry, *value* is the *default*.

Present-p is *true* if an entry is found; otherwise, it is *false*.

`setf` may be used with `gethash` to modify the *value* associated with a given *key*, or to add a new entry. When a `gethash` *form* is used as a `setf` *place*, any *default* which is supplied is evaluated according to normal left-to-right evaluation rules, but its *value* is ignored.

Examples:

```
(setq table (make-hash-table)) → #<HASH-TABLE EQL 0/120 32206334>
(gethash 1 table) → NIL, false
(gethash 1 table 2) → 2, false
(setf (gethash 1 table) "one") → "one"
(setf (gethash 2 table "two") "two") → "two"
(gethash 1 table) → "one", true
(gethash 2 table) → "two", true
(gethash nil table) → NIL, false
(setf (gethash nil table) nil) → NIL
(gethash nil table) → NIL, true
(defvar *counters* (make-hash-table)) → *COUNTERS*
(gethash 'foo *counters*) → NIL, false
(gethash 'foo *counters* 0) → 0, false
```

```
(defmacro how-many (obj) `(values (gethash ,obj *counters* 0))) → HOW-MANY
(defun count-it (obj) (incf (how-many obj))) → COUNT-IT
(dolist (x '(bar foo foo bar bar baz)) (count-it x))
(how-many 'foo) → 2
(how-many 'bar) → 3
(how-many 'quux) → 0
```

See Also:

`remhash`

Notes:

The *secondary value*, *present-p*, can be used to distinguish the absence of an entry from the presence of an entry that has a value of *default*.

remhash

Function

Syntax:

`remhash key hash-table` → *generalized-boolean*

Arguments and Values:

key—an *object*.

hash-table—a *hash table*.

generalized-boolean—a *generalized boolean*.

Description:

Removes the entry for *key* in *hash-table*, if any. Returns *true* if there was such an entry, or *false* otherwise.

Examples:

```
(setq table (make-hash-table)) → #<HASH-TABLE EQL 0/120 32115666>
(setf (gethash 100 table) "C") → "C"
(gethash 100 table) → "C", true
(remhash 100 table) → true
(gethash 100 table) → NIL, false
(remhash 100 table) → false
```

Side Effects:

The *hash-table* is modified.

maphash

maphash

Function

Syntax:

`maphash function hash-table` → nil

Arguments and Values:

function—a *designator* for a *function* of two *arguments*, the *key* and the *value*.

hash-table—a *hash table*.

Description:

Iterates over all entries in the *hash-table*. For each entry, the *function* is called with two *arguments*—the *key* and the *value* of that entry.

The consequences are unspecified if any attempt is made to add or remove an entry from the *hash-table* while a **maphash** is in progress, with two exceptions: the *function* can use **setf** of **gethash** to change the *value* part of the entry currently being processed, or it can use **remhash** to remove that entry.

Examples:

```
(setq table (make-hash-table)) → #<HASH-TABLE EQL 0/120 32304110>
(dotimes (i 10) (setf (gethash i table) i)) → NIL
(let ((sum-of-squares 0))
  (maphash #'(lambda (key val)
    (let ((square (* val val)))
      (incf sum-of-squares square)
      (setf (gethash key table) square)))
    table)
  sum-of-squares) → 285
(hash-table-count table) → 10
(maphash #'(lambda (key val)
  (when (oddp val) (remhash key table)))
  table) → NIL
(hash-table-count table) → 5
(maphash #'(lambda (k v) (print (list k v))) table)
(0 0)
(8 64)
(2 4)
(6 36)
(4 16)
→ NIL
```

Side Effects:

None, other than any which might be done by the *function*.

See Also:

loop, with-hash-table-iterator, Section 3.6 (Traversal Rules and Side Effects)

with-hash-table-iterator

Macro

Syntax:

with-hash-table-iterator (*name hash-table*) {*declaration*}* {*form*}* → {*result*}*

Arguments and Values:

name—a name suitable for the first argument to **macrolet**.

hash-table—a *form*, evaluated once, that should produce a *hash table*.

declaration—a **declare** *expression*; not evaluated.

forms—an *implicit progn*.

results—the *values* returned by *forms*.

Description:

Within the lexical scope of the body, *name* is defined via **macrolet** such that successive invocations of (*name*) return the items, one by one, from the *hash table* that is obtained by evaluating *hash-table* only once.

An invocation (*name*) returns three values as follows:

1. A *generalized boolean* that is *true* if an entry is returned.
2. The key from the *hash-table* entry.
3. The value from the *hash-table* entry.

After all entries have been returned by successive invocations of (*name*), then only one value is returned, namely **nil**.

It is unspecified what happens if any of the implicit interior state of an iteration is returned outside the dynamic extent of the **with-hash-table-iterator** *form* such as by returning some *closure* over the invocation *form*.

Any number of invocations of **with-hash-table-iterator** can be nested, and the body of the innermost one can invoke all of the locally *established macros*, provided all of those *macros* have *distinct* names.

Examples:

The following function should return **t** on any *hash table*, and signal an error if the usage of **with-hash-table-iterator** does not agree with the corresponding usage of **maphash**.

```
(defun test-hash-table-iterator (hash-table)
  (let ((all-entries '())
        (generated-entries '())
        (unique (list nil)))
    (maphash #'(lambda (key value) (push (list key value) all-entries))
              hash-table)
    (with-hash-table-iterator (generator-fn hash-table)
      (loop
        (multiple-value-bind (more? key value) (generator-fn)
          (unless more? (return))
          (unless (eql value (gethash key hash-table unique))
            (error "Key ~S not found for value ~S" key value))
          (push (list key value) generated-entries))))
    (unless (= (length all-entries)
               (length generated-entries)
               (length (union all-entries generated-entries
                              :key #'car :test (hash-table-test hash-table))))
      (error "Generated entries and Maphash entries don't correspond"))
    t))
```

The following could be an acceptable definition of **maphash**, implemented by **with-hash-table-iterator**.

```
(defun maphash (function hash-table)
  (with-hash-table-iterator (next-entry hash-table)
    (loop (multiple-value-bind (more key value) (next-entry)
          (unless more (return nil))
          (funcall function key value)))))
```

Exceptional Situations:

The consequences are undefined if the local function named *name established* by **with-hash-table-iterator** is called after it has returned *false* as its *primary value*.

See Also:

Section 3.6 (Traversal Rules and Side Effects)

clrhash

Function

Syntax:

`clrhash hash-table` \rightarrow *hash-table*

Arguments and Values:

hash-table—a *hash table*.

Description:

Removes all entries from *hash-table*, and then returns that empty *hash table*.

Examples:

```
(setq table (make-hash-table)) → #<HASH-TABLE EQL 0/120 32004073>
(dotimes (i 100) (setf (gethash i table) (format nil "~R" i))) → NIL
(hash-table-count table) → 100
(gethash 57 table) → "fifty-seven", true
(clearhash table) → #<HASH-TABLE EQL 0/120 32004073>
(hash-table-count table) → 0
(gethash 57 table) → NIL, false
```

Side Effects:

The *hash-table* is modified.

sxhash

Function

Syntax:

`sxhash object` → *hash-code*

Arguments and Values:

object—an *object*.

hash-code—a non-negative *fixnum*.

Description:

`sxhash` returns a hash code for *object*.

The manner in which the hash code is computed is *implementation-dependent*, but subject to certain constraints:

1. (`equal x y`) implies (`= (sxhash x) (sxhash y)`).
2. For any two *objects*, *x* and *y*, both of which are *bit vectors*, *characters*, *conses*, *numbers*, *pathnames*, *strings*, or *symbols*, and which are *similar*, (`sxhash x`) and (`sxhash y`) yield the same mathematical value even if *x* and *y* exist in different *Lisp images* of the same *implementation*. See Section 3.2.4 (Literal Objects in Compiled Files).
3. The *hash-code* for an *object* is always the *same* within a single *session* provided that the *object* is not visibly modified with regard to the equivalence test **equal**. See Section 18.1.2 (Modifying Hash Table Keys).

sxhash

4. The *hash-code* is intended for hashing. This places no verifiable constraint on a *conforming implementation*, but the intent is that an *implementation* should make a good-faith effort to produce *hash-codes* that are well distributed within the range of non-negative *fixnums*.
5. Computation of the *hash-code* must terminate, even if the *object* contains circularities.

Examples:

```
(= (sxhash (list 'list "ab")) (sxhash (list 'list "ab"))) → true
(= (sxhash "a") (sxhash (make-string 1 :initial-element #\a))) → true
(let ((r (make-random-state)))
  (= (sxhash r) (sxhash (make-random-state r))))
→ implementation-dependent
```

Affected By:

The *implementation*.

Notes:

Many common hashing needs are satisfied by **make-hash-table** and the related functions on *hash tables*. **sxhash** is intended for use where the pre-defined abstractions are insufficient. Its main intent is to allow the user a convenient means of implementing more complicated hashing paradigms than are provided through *hash tables*.

The hash codes returned by **sxhash** are not necessarily related to any hashing strategy used by any other *function* in Common Lisp.

For *objects* of *types* that **equal** compares with **eq**, item 3 requires that the *hash-code* be based on some immutable quality of the identity of the object. Another legitimate implementation technique would be to have **sxhash** assign (and cache) a random hash code for these *objects*, since there is no requirement that *similar* but non-**eq** objects have the same hash code.

Although *similarity* is defined for *symbols* in terms of both the *symbol's name* and the *packages* in which the *symbol* is *accessible*, item 3 disallows using *package* information to compute the hash code, since changes to the package status of a symbol are not visible to **equal**.
